

# Project Report: Predictive Maintenance using Sensor Data

By Promise Ekpo

## 1. Problem Statement & Business Context

**Problem:** Predict generator failures in wind-turbine equipment using sensor data so that components can be repaired before catastrophic failure, lowering maintenance and replacement costs.

**Business context:** ReneWind (a wind-energy company) provided a ciphered sensor data set. Predictive maintenance reduces expensive replacements by forecasting failures; repairing is much cheaper than replacing, and inspections are cheaper than repairs. The cost structure motivates maximizing detection of true failures (minimizing false negatives).

**Source:** Dataset as well as project Git hub repo is available from the below link:  
[promwizy911/Predictive-Maintenance-using-Sensor-Data](https://github.com/promwizy911/Predictive-Maintenance-using-Sensor-Data)

## 2. Dataset Description

**Source and structure:** The data is a transformed/ciphered sensor data set (Train.csv and Test.csv). Each set contains 40 predictor variables (V1..V40) and a binary target (Target: 1=failure, 0=no failure). Train set: 20,000 rows; Test set: 5,000 rows. ↴

**Key statistics and imbalance:** The target distribution in the training set is 18,890 of '0' and 1,110 of '1' (~5.6% failures), so the data set is highly imbalanced. Summary statistics (means, std, ranges) and uni-variate plots were used during EDA. ↴

**Missing data & challenges:** Two features, V1 and V2, contained 18 missing values each; other columns were complete. Features are numerical (float64) except Target (int). Correlations between features are generally weak with a few moderate relationships identified. ↴

**Limitations / biases:** Data set is ciphered (feature names obfuscated) so domain-meaning interpretation of individual predictors is not possible from the file alone. The imbalance (few failure cases) requires careful sampling or algorithmic treatment to avoid biased models.

## 3. Data Preprocessing & Feature Engineering

Splits used in development: An initial stratified split created training and validation sets (X\_train: 14,000; X\_val: 6,000). For the final pipeline an 80/20 split was also illustrated (X\_train: 4,000; X\_test: 1,000) when building the production pipeline. ↴[filecite](#) ↴[turn1file17](#) ↴[turn1file2](#) ↴

**Missing-value handling:** KNNImputer (k=5) was used to impute missing values in V1 and V2 on training and validation sets to avoid leakage. After imputation there were no missing values left.

**Sampling / imbalance handling:** The team experimented with (a) training on original data, (b) oversampling minority class using SMOTE, and (c) under sampling using RandomUnderSampler. Under-sampling produced models with better recall and less over fitting for some algorithms; oversampling improved recall but increased over fitting in others.

**Feature engineering:** The data set was already transformed/ciphered; EDA identified feature relationships (e.g., strong correlations among V15–V7 and negative correlation between V7 and V18) used to guide modelling. The final pipeline used ColumnTransformer + KNN imputer to wrap preprocessing. No explicit domain-derived features are added in the notebook.

## 4. Model Architecture & Approach

**Model family:** The project focuses on classical supervised classifiers (tree-based ensembles and linear models). Models trained and compared include DecisionTree, Bagging, RandomForest, GradientBoosting (GBM), AdaBoost, XGBoost, RidgeClassifier, LogisticRegression (and Lasso).

Why GBM was chosen as final approach: Gradient Boosting (GBM) showed strong recall performance across sampling strategies and, when tuned and combined with under-sampling, produced the best generalization (high recall with relatively low over-fitting). The choice is empirical — GBM provided the best trade-off for maximizing recall (the priority metric). ☑

**Conceptual workflow (preprocessing → sampling → model):**

1. Read Train.csv / Test.csv → 2. EDA & missing-value detection → 3. Impute V1,V2 with KNNImputer → 4. Optionally resample data (SMOTE or undersample) → 5. Train candidate models with stratified K-Fold CV optimizing Recall → 6. Hyperparameter tuning (Randomized / Grid search) → 7. Final pipeline assembly with preprocessing + undersampling + tuned GBM → 8. Evaluate on held-out test set. ☑

## 5. Training Strategy

**Train / validation approach:** A stratified K-Fold ( $k=5$ ) cross-validation was used to assess models on training data. An explicit train/validation split was also used (70% train, 30% validation) for rapid evaluation. For final pipeline testing an independent hold-out test set was used.

**Scoring metric and loss:** Recall (sensitivity) was selected as the primary scorer because minimizing false negatives (missed failures) has the largest business cost. Scorer used in cross-validation and hyperparameter search = recall\_score.

**Hyperparameter tuning:** RandomizedSearchCV and GridSearchCV were used for GBM, RandomForest and XGBoost across original, over-sampled and under-sampled data sets. Final tuned hyperparameters for GBM (under-sampled best): learning\_rate=0.15, max\_features=0.67, n\_estimators=87, subsample=0.2. Random and grid searches reported cross-validated recall scores for candidate parameter sets.

Notes on optimizers / loss: The models used (tree ensembles and XGBoost) rely on their internal boosting objectives (log-loss / deviance for GBM), and XGBoost used its standard objectives (eval\_metric='logloss' shown). The notebook did not apply neural-network optimizers. ☑

## 6. Model Evaluation

**Primary metrics reported:** Accuracy, Recall, Precision, and F1. Recall was prioritized. A confusion-matrix helper and a model\_performance function were used to compute and tabulate these metrics. ☑

**Highlights of results (representative numbers from the notebook):**

- Best baseline XGBoost on original data: validation recall  $\approx 0.825$ . ☑
- GBM tuned on original data (GridSearch/Randomized): validation recall  $\approx 0.79\text{--}0.80$  but with overfitting observed. ☑
- GBM trained on oversampled data achieved CV recall  $\approx 0.97$  (training/oversampled CV) but showed signs of overfitting; validation recall lower ( $\sim 0.81\text{--}0.88$  depending on tuning).
- GBM trained and tuned on undersampled data produced strong recall with reduced overfitting; example tuned undersampled GBM produced CV recall  $\approx 0.88\text{--}0.90$  and validation recall  $\approx 0.88\text{--}0.89$ . This was judged the best generalizing model. ☑

**Final pipeline evaluation on test split:** The assembled pipeline (ColumnTransformer KNN imputer  $\rightarrow$  undersampling transformer  $\rightarrow$  GBM with tuned parameters) evaluated on the held-out test set reported: Accuracy=0.964, Recall=0.625, Precision=0.700, F1=0.660. (These are the notebook's reported test metrics for the final pipeline run.) ☑filecite☐turn1file2☐

**Comparison with baselines:** The project compared multiple algorithms (DecisionTree, Bagging, RandomForest, GBM, XGBoost, logistic/ridge/lasso). GBM (with appropriate sampling/tuning) and XGBoost were consistently top performers for recall. Undersampling + tuned GBM offered better generalization versus over-sampled variants that tended to over-fit.

## 7. Deployment / MLOps (If Applicable)

**Reproducible pipeline:** The notebook assembles an sklearn Pipeline that encapsulates preprocessing (KNN imputation via ColumnTransformer), an undersampling FunctionTransformer, and the final tuned GradientBoostingClassifier. This pipeline structure supports serialization (e.g., via joblib/pickle) and is deployment-friendly.

**CI/CD & monitoring:** The project demonstrated pipeline creation and evaluation but did not include explicit CI/CD scripts, model registry steps, or runtime monitoring dashboards. Those would be recommended next steps but were not done in the project.

**Scalability considerations:** Tree ensembles are CPU-bound at training time but fast at inference and can be served in a standard model-serving environment (REST API, batch jobs). The use of a

single pipeline simplifies reproducible preprocessing during serving. The project itself does not include auto-scaling or production monitoring code. ☐

## 8. Key Insights & Business Impact

**Top predictors:** The most important features in the final model are V18 (top), V36, V15, V3, and V26 — these features correlate with failure risk and are actionable triggers for maintenance.

☒filecite☒turn1file2☒

**Actionable rule-of-thumb (from EDA & model interpretation):** Conditions such as  $V15 > -5$  combined with  $V7 > 0$ , or  $V18$  dropping below 0 with  $V36$  below thresholds, significantly raise failure probability. The project suggested immediate maintenance when  $V15$  increases above  $-5$ . These rules provide tangible operational triggers for inspection/repair.

**Expected business impact:** Using the model to flag likely failures allows pre-emptive repairs (cheaper than replacements), reducing downtime and lowering maintenance costs. The focus on Recall aligns model objectives with the business cost structure described. ☐

## 9. Final Conclusion

**Learning and challenges:** The project shows that:

- (a) careful treatment of class imbalance is critical.
- (b) ensemble models (GBM/XGBoost) perform well for failure detection.
- (c) oversampling can increase recall but risks over-fitting. Missing values were minimal and handled with KNN imputation. Feature names were ciphered, limiting domain interpretability.

**Future scope and improvements:**

- Collect richer labeled failure examples to reduce imbalance and increase confidence in rare-event prediction.
- Add domain-meaningful feature engineering if original sensor labels become available (e.g., aggregations over time, rolling statistics).
- Implement model versioning, automated CI/CD deployment, and production monitoring (drift detection, alerting).
- Consider cost-sensitive learning or custom cost function optimization to directly minimize expected maintenance + replacement costs. (Not implemented in the provided notebook.)